

# **ECE 3640 - Discrete-Time Signals and Systems**

## **Matlab: Basic Knowledge and Skills**

Jake Gunther

Spring 2015



Department of Electrical & Computer Engineering

# outline

- signal input and output using Matlab's media read and write functions
  - audio: wavread/wavwrite and audioread/audiowrite
  - images: imread/imwrite
  - videos: videoReader/videoWriter
- multidimensional arrays in Matlab
- Matlab low-level file I/O: fopen, fclose, fread, fwrite (interface to C)
- `plot` function
  - plotting and signal (one independent variable)
  - embellishments
  - handle graphics
- `image` and `imagesc` function
  - plotting and image (two independent variables)
  - embellishments
  - handle graphics
- Fourier transforms via `fft` and spectral plots
- `spectrogram` function and time-frequency plots

help in Matlab

# getting help

- `help` function - brings up information about the function
- `lookfor` keyword - searches for that keyword throughout the help system
- `doc` - brings up the documentation system
- `helpdesk`
- `www.mathworks.com`
- ask a human
- ask the Internet (i.e. Google)

# example

```
>> help relop
Relational operators.
< > Relational operators.
    The six relational operators are <, <=, >, >=, ==, and ~=.
    A < B does element by element comparisons between A and B ...

& Element-wise Logical AND.
    A & B is a matrix whose elements are logical 1 (TRUE) where both A
    and B have non-zero elements, and logical 0 (FALSE) where either has ...

&& Short-Circuit Logical AND.
    A && B is a scalar value that is the logical AND of scalar A and B.
    This is a "short-circuit" operation in that MATLAB evaluates B only ...

| Element-wise Logical OR.
    A | B is a matrix whose elements are logical 1 (TRUE) where either
    A or B has a non-zero element, and logical 0 (FALSE) where both have ...

|| Short-Circuit Logical OR.
    A || B is a scalar value that is the logical OR of scalar A and B.
    This is a "short-circuit" operation in that MATLAB evaluates B only ...

~ Logical complement (NOT).
    ~A is a matrix whose elements are logical 1 (TRUE) where A has zero
    elements, and logical 0 (FALSE) where A has non-zero elements.

xor Exclusive OR.
    xor(A,B) is logical 1 (TRUE) where either A or B, but not both, is
    non-zero. See XOR.
```



# useful Matlab commands

- `pwd` - show the present working directory
- `dir` and `ls` - lists all files in the current directory
- `cd ./bob` and `chdir ./bob` - changes to the bob subdirectory
- `what` - lists all m-files in the current directory
- `which test` display path to `test.m`
- `type test` - display `test.m` in command window
- `delete test` - deletes file `test.m`
- `who` and `whos` - displays all variables in the workspace
- `save` - saves all the variables in the workspace
- `load` - loads variables into the workspace
- `clear all` - clears all variables from the workspace
- `clear x` - clears only the variable `x` from the workspace

# signal input and output in Matlab

# different kinds of signals

signal	dimension	number channels
audio (monaural)	1	1
audio (stereo)	1	2
image (grayscale)	2	1
image (RGB color)	2	3
video (grayscale)	3	1
video (RGB color)	3	3






# media: read audio signals

```
1 ai = audioinfo('../siren.wav');  
2 disp(ai); % display audio header structure  
3 [x,fs] = audioread('../siren.wav',[1 10]*ai.SampleRate); % read audio file  
4 soundsc(x,fs); % play sound to speaker  
5 t = [0:length(x)-1]/fs;  
6 plot(t,x);  
7 xlabel('Time [seconds]');  
8 ylabel('Amplitude');  
9 title('Signal read from the file "siren.wav"');  
10 print -dpng siren.png
```

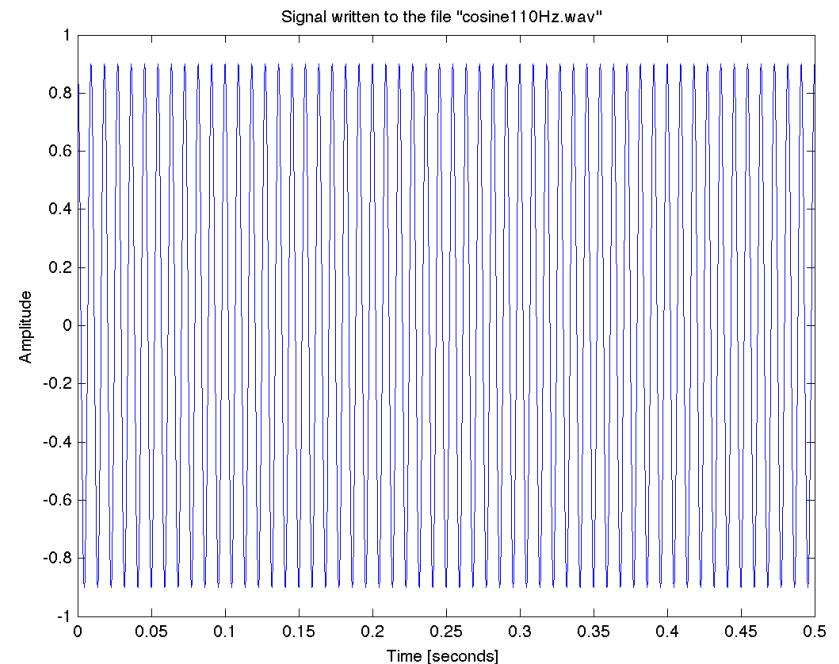
Filename: 'siren.wav'  
CompressionMethod: 'Uncompressed'  
NumChannels: 1  
SampleRate: 44100  
TotalSamples: 2182685  
Duration: 49.4940  
Title: []  
Comment: []  
Artist: []  
BitsPerSample: 16



# media: write audio signals

```
1 fs = 8000; % sample rate [samples/second]
2 t = [0:fs*0.5-1]/fs;
3 x = 0.9*cos(2*pi*110*t);
4 audiowrite('cosine110Hz.wav',x,fs);
5 ai = audioinfo('cosine110Hz.wav');
6 disp(ai);
7 soundsc(x,fs); % play sound to speaker
8 plot(t,x);
9 xlabel('Time [seconds]');
10 ylabel('Amplitude');
11 title('Signal written to the file "cosine110Hz.wav"');
12 print -dpng cosine110Hz.png
```

Filename: 'cosine110Hz.wav'  
CompressionMethod: 'Uncompressed'  
NumChannels: 1  
SampleRate: 8000  
TotalSamples: 4000  
Duration: 0.5000  
Title: []  
Comment: []  
Artist: []  
BitsPerSample: 16





# media: read images

```
1 clear all; clf;
2 %imfinfo('cameraman.tif')
3 x = imread('cameraman.tif');
4 subplot(121);
5 image(x); % show the image
6 axis image; % make the pixels square
7 title('Unscaled - image');
8 subplot(122);
9 imagesc(x); % show scaled image
10 axis image; % make the pixels square
11 title('Scaled - imagesc');
12 colormap gray; % this is a grayscale image
13 print -dpng cameraman_scaling.png
14 [nrows,ncols] = size(x);
15 fprintf('This image is %d X %d pixels.\n\n',nrows,ncols);
16 whos
```

This image is 256 X 256 pixels.

Name	Size	Bytes	Class	Attributes
ncols	1x1	8	double	
nrows	1x1	8	double	
x	256x256	65536	uint8	




## media: write images

```
1 clear all; clf;
2 x = dftmtx(16);
3 x = [real(x) imag(x)];
4 imwrite(x, 'dftmtx.png', 'PNG');
5 imagesc(x);
6 axis image;
7 colormap(jet);
8 print -dpng dftmtx_matlab.png
9 [nrows,ncols] = size(x);
10 fprintf('This image is %d X %d pixels.\n\n',nrows,ncols);
11 whos
```

This image is 16 X 32 pixels.

Name	Size	Bytes	Class	Attributes
ncols	1x1	8	double	
nrows	1x1	8	double	
x	16x32	4096	double	

(in the image file) → 



(in the Matlab figure window)


# media: read color images

```
1 clear all; clf;
2 x = imread('coloredChips.png');
3 r = x; r(:,:, [2,3]) = 0;
4 g = x; g(:,:, [1,3]) = 0;
5 b = x; b(:,:, [1,2]) = 0;
6 image([x r; g b]);
7 axis image; % make the pixels square
8 print -dpng coloredChips_color_planes.png;
9 [nrows,ncols,nrgb] = size(x);
10 fprintf('This image is %d X %d pixels X %d color planes.\n\n',...
11         nrows,ncols,nrgb);
12 whos
```



This image is 391 X 518 pixels X 3 color planes.

Name	Size	Bytes	Class	Attributes
b	391x518x3	607614	uint8	
g	391x518x3	607614	uint8	
ncols	1x1	8	double	
nrgb	1x1	8	double	
nrows	1x1	8	double	
r	391x518x3	607614	uint8	
x	391x518x3	607614	uint8	




# media: write color images

```
1 clear all; clf;
2 x = dftmtx(16);
3 r = abs(real(x));
4 g = abs(imag(x));
5 b = abs(real(x));
6 x = cat(3,r,g,b); % concatenate along the third dimension
7 imwrite(x,'dftmtx_color.png','PNG');
8 image(x);
9 axis image; % make the pixels square
10 print -dpng dftmtx_color_matlab.png;
11 [nrows,ncols,nrgb] = size(x);
12 fprintf('This image is %d X %d pixels X %d color planes.\n\n',...
13         nrows,ncols,nrgb);
14 whos
```

his image is 16 X 16 pixels X 3 color planes.

(in the image file) → 


Name	Size	Bytes	Class	Attributes
b	16x16	2048	double	
g	16x16	2048	double	
ncols	1x1	8	double	
nrgb	1x1	8	double	
nrows	1x1	8	double	
r	16x16	2048	double	
x	16x16x3	6144	double	



(in the Matlab figure window)

## media: videos

- grayscale videos are 3-dimensional: (row,col,frame)
- color videos are 4-dimensional: (row,col,rgb,frame)
- video files can be very large
- read and write one frame at a time




## media: read videos

```
1 clear all; clf;
2 obj = VideoReader('xylophone.mp4', 'tag', 'myreader1');
3 x = read(obj); % read in all the frames
4 fprintf('video x: '); disp(size(x));
5 y = read(obj,20); % read in the i-th frame
6 fprintf('frame y: '); disp(size(y));
7 % read in the frames one at a time
8 nframes = obj.NumberOfFrames;
9 for i=1:nframes
10     y = read(obj,i);
11     if(i==1)
12         imhan = image(y); axis image; drawnow;
13     else
14         set(imhan,'CData',y); drawnow;
15     end
16 end
```

video x:     240     320     3     141

frame y:     240     320     3



(first frame of the video)



## media: write videos

```
1 clear all; clf;
2 obj = VideoWriter('wave.mp4','MPEG-4'); % create video file
3 open(obj); % open video file
4 fs = 30; % sample rate [frames/second]
5 t = [0:60]/fs; y = [0:31]/50; x = [0:63]/50; [Y,X]=meshgrid(y,x);
6 for i=1:length(t)
7     Zc = cos(2*pi*(t(i)*1 - X*1 - Y*0.75)).^2;
8     Zs = sin(2*pi*(t(i)*1 - X*1 - Y*0.75)).^2;
9     Z = cat(3,Zc,Zs,flipplr(Zc)); % make a video frame = RGB image
10    writeVideo(obj,Z); % write the video frame
11    if(i==1)
12        imhan = image(Z); drawnow;
13    else
14        set(imhan,'CData',Z); drawnow;
15    end
16 end
17 close(obj); % close video file
18 % See also "getframe" to make a video of the Matlab figure window.
```

- watch the video

## **multidimensional arrays in Matlab**

# MD arrays: terminology

- 1D array can be a row vector or a column vector

```
1 x = zeros(1,4)
2 y = ones(3,1)
```


```
x =
    0     0     0     0

y =
    1
    1
    1
```

- 2D array has both rows and columns

```
1 x = rand(2,3)
```

```
x =
    0.2621    0.7549    0.4424
    0.0445    0.2428    0.6878
```



(3D array representing a color image)

- 3D arrays have rows, columns and pages/slabs (see picture)
- 4D arrays have rows, columns, pages, and books
- 5D arrays have rows, columns, pages, books and shelves
- 6D arrays have rows, columns, pages, books, shelves and aisles

# creating MD arrays: built-in functions

```
1 x = floor(10*rand(2,5,6)) % create a random 4D array
2 n = ndims(x) % return the number of dimensions
3 s = size(x) % return the size of x in each dimension
```

```
x(:,:,1) =
    9    7    1    4    6
    2    2    6    4    7

x(:,:,2) =
    3    4    8    6    5
    6    8    2    5    8

x(:,:,3) =
    2    1    6    6    6
    3    9    4    5    5


x(:,:,4) =
    7    9    1    0    4
    5    2    1    4    3

x(:,:,5) =
    7    7    9    1    0
    6    9    1    6    5

x(:,:,6) =
    5    4    6    5    1
    8    3    7    3    5

n =
    3

s =
    2    5    6
```



- functions `rand`, `randn`, `ones`, `zeros` produce MD arrays
- can do 4D `ones([2 3 4 5])` and 5D `zeros([4 2 20 75 8])` and so on

# creating MD arrays: concatenation

```

1 a = floor(10*rand(3,4));
2 b = ones(3,4);
3 c = zeros(3,4);
4 d = round(10*randn(3,4));
5 x = cat(3,a,b,c,d)
6 n = ndims(x)           % return the number of dimensions
7 s = size(x)            % return the size of x in each dimension

```

```

x(:,:,1) =
     0     4     6     2
     1     1     1     9
     7     3     7     2

x(:,:,2) =
     1     1     1     1
     1     1     1     1
     1     1     1     1


x(:,:,3) =
     0     0     0     0
     0     0     0     0
     0     0     0     0

x(:,:,4) =
     2    -6     3     8
    -15     4    -4     1
     -7     9     8     6

n =
     3

s =
     3     4     4

```



# creating MD arrays: replication

```
1 x = repmat(eye(3), [1,1,4])
2 n = ndims(x)           % return the number of dimensions
3 s = size(x)            % return the size of x in each dimension
```

```
x(:,:,1) =
    1     0     0
    0     1     0
    0     0     1
x(:,:,2) =
    1     0     0
    0     1     0
    0     0     1
x(:,:,3) =
    1     0     0
    0     1     0
    0     0     1
x(:,:,4) =
    1     0     0
    0     1     0
    0     0     1
n =
    3
s =
    3     3     4
```

- can do 4D `repmat(eye(3), [1 1 4 6])`
- can do 5D `repmat(eye(3), [1 1 4 100 35])` and so on

# creating MD arrays: reshaping

```
1 x = reshape([1:48],[2 3 2 4])
2 n = ndims(x)           % return the number of dimensions
3 s = size(x)            % return the size of x in each dimension
```

```
x(:,:,1,1) =
     1     3     5
     2     4     6
x(:,:,2,1) =
     7     9    11
    10    12
x(:,:,1,2) =
    13    15    17
    14    16    18
x(:,:,2,2) =
    19    21    23
    20    22    24
x(:,:,1,3) =
    25    27    29
    26    28    30
x(:,:,2,3) =
    31    33    35
    32    34    36
x(:,:,1,4) =
    37    39    41
    38    40    42
x(:,:,2,4) =
    43    45    47
    44    46    48
n =
     4
s =
     2     3     2     4
```

- this example illustrates how Matlab's MD arrays are laid out in memory: the leftmost subscript varies fastest as elements are accessed in memory
- this is opposite to the way that MD arrays are laid out in C

## creating MD arrays: reading in media

```
1 [v,fs] = audioread('axe-head.wav'); % mono wav file
2 sv = size(v)
3 [w,fs] = audioread('music.mp3'); % stereo wav file
4 sw = size(w)
5 x = imread('cameraman.tif'); % grayscale image
6 sx = size(x)
7 y = imread('coloredChips.png'); % color image
8 sy = size(y)
9 obj = VideoReader('xylophone.mp4','tag','myreader1');
10 z = read(obj); % color video
11 sz = size(z)
```

```
sv =
    187165         1    <= mono audio

sw =
    2087424         2    <= stereo audio

sx =
    256    256         <= grayscale image

sy =
    391    518         3    <= color image

sz =
    240    320         3    141    <= color video
```




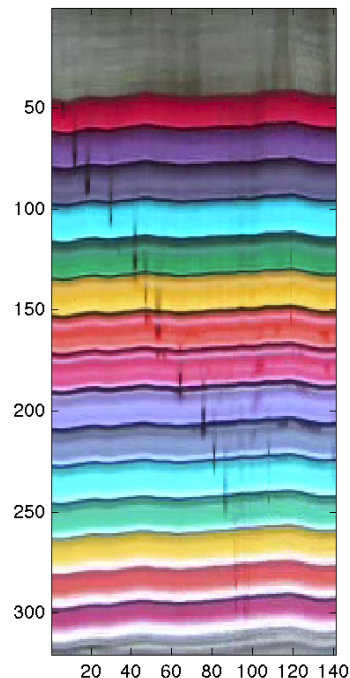
## accessing MD array elements

- get and set elements of MD arrays by reference to specific element
  - 1D: `x(30000)`
  - 2D: `x(30000, 2)` or `x(100, 200)`
  - 3D: `x(100, 200, 2)`
  - 4D: `x(100, 200, 2, 75)` or `x(i, j, k, l)`
- slicing using colon (:) to get and set
  - 1D: `x(1:14000)` or `x(14000:end)` or `x(1:2:end)`
  - 2D: `x(30:35, 70:75)` or `x([30, 23, 99], [10:2:20, 70:5:end])`
  - 3D: `x(:, :, 3)` accesses the blue plane in an RGB color image
  - 4D: `x(:, :, :, 121)` accesses the 121st frame in a color video, returns a color image
  - 4D: `x(100, 200, 3, :)` accesses the (100,200)th pixel, blue color plane over all the frames in a video
- `x(:)` vectorizes the object (returns a column vector)

# squeezing and permuting by example

```
1 obj = VideoReader('xylophone.mp4','tag','myreader1');
2 z = read(obj); % color video
3 sz = size(z)
4 r = round(sz(1)/2)
5 y = z(r,:,:,);
6 sy = size(y)
7 x = squeeze(y); % squeeze removes singleton dimensions
8 sx = size(x)
9 w = permute(x,[1 3 2]); % permutes the dimensions, need RGB in 3rd
10                        % dimension of image (see also ipermute)
11 sw = size(w)
12 imagesc(w); axis image;
```

```
sz =
    240    320     3    141
r =
    120
sy =
     1    320     3    141
sx =
    320     3    141
sw =
    320    141     3
```



# operations on MD arrays

- arithmetic: `+`, `-`, `.*`, `./`, `.^`, etc.
- functions that do not change dimension: `round`, `mod`, `log`, `exp`, `abs`, `cos`, `tanh`, etc.
- functions that reduce dimension by 1: `sum`, `mean`, `std`, `var`, `min`, `max`, etc.
- example: `max(max(max(x)))` finds the largest element in a 3D array
- example: `mean(mean(y))` finds the average R, G and B values in a color image

## low-level file I/O in Matlab

## example file I/O

```
1 fid = fopen('data.bin','rb');
2 if(fid == -1) fprintf('ERROR: Could not open file.\n'); end
3 [x,cnt] = fread(fid, 5, 'int'); % read in 5 integers
4 [y,cnt] = fread(fid, 10, 'float'); % read in 10 floats
5 [z,cnt] = fread(fid,inf, 'short'); % read to the end of the file
6 fclose(fid);
7
8 y = rand(10,1); % 10 double-precision uniform [0,1] random numbers
9 fid = fopen('data.bin','wb');
10 cnt = fwrite(fid,y, 'float'); % write values in y array as floats
11 fclose(fid);
```

- you have to know the file format to read it correctly
- use `fclose all` to close all open files
- other functions available: `fseek`, `frewind`, `ftell`, `feof`, `ferror`, `fscanf`, `fgetl`, `fgets`, etc.

the `plot` command

## using plot

- `plot` plots points connected by line segments
- points are provided to `plot` using equal length vectors:
  - `plot(x,y)` - plots `y` versus `x` by connecting the points with line segments
  - `plot(y)` - plots `y` versus the index `[1:length(y)]`
- plotting multiple data sets on the same axis done in two ways
  - `plot(x1,y1,x2,y2,x3,y3)` plots 3 lines using different colors for each
  - add lines to an existing plot using `hold on` and `hold off`

```
1 plot(x1,y1,'Color',[0 0 1]);    % create a plot
2 hold on;    % add subsequent plots to this axis
3 plot(x2,y2,'Color',[0 1 0]);    % add a plot
4 plot(x3,y3,'Color',[1 0 0]);    % add a plot
5 hold off; % stop adding plots to this axis
```

does the same thing as

```
1 plot(x1,y1,'b'); hold on;
2 plot(x2,y2,'g');
3 plot(x3,y3,'r'); hold off;
```

## predefined colors, markers and line types

b	blue	.	point	-	solid
g	green	o	circle	:	dotted
r	red	x	x-mark	-.	dashdot
c	cyan	+	plus	--	dashed
m	magenta	*	star	(none)	no line
y	yellow	s	square		
k	black	d	diamond		
w	white	v	triangle (down)		
		^	triangle (up)		
		<	triangle (left)		
		>	triangle (right)		
		p	pentagram		
		h	hexagram		

- `plot(x, y, 'm*--')` plots magenta line, “\*” marker, and dashed line
- `plot(x, y, 'k', x, y, 'r+')` (what does this do?)



# handle graphics

- a plot has many different attributes that can be examined/changed using `get/set`

```
1 x = [0:8]*pi/8; y = sin(x);  
2 han = plot(x,y)  
3 get(han);
```

```
Color: [0 0 1]  
LineStyle: '-'  
LineWidth: 0.5000  
Marker: 'none'  
MarkerSize: 6  
MarkerEdgeColor: 'auto'  
MarkerFaceColor: 'none'  
XData: [1x21 double]  
YData: [1x21 double]  
Type: 'line'  
Visible: 'on'  
Parent: 173.0143
```

- to see the attributes of a graphics handle use `get(han)`
- to see the attributes of a graphics handle that the user can set use `set(han)`
- to set the attributes use `set(han, 'attribute', value)`

- then do

```
1 set(han, 'LineWidth', 2, 'Marker', '^', 'MarkerSize', 12);  
2 set(han, 'MarkerFaceColor', [1 0 1]);
```

- each plot line has its own attributes and can be set independently

# figures

- plot lines are children of an axes objects
- axes are children of figure objects
- figures and axes have many properties that can be set by the user
- `gcf` returns a handle to the current figure
- `get(gcf)` lists the attributes of a figure
- below is only a partial list

```
Color = [0.8 0.8 0.8]  
CurrentAxes = [173.014]  
PaperUnits = inches  
PaperOrientation = portrait  
PaperPosition = [0.25 2.5 8 6]  
PaperPositionMode = manual
```

```
PaperSize = [8.5 11]  
PaperType = usletter  
Position = [1201 440 560 420]  
Children = [173.014]  
Parent = [0]  
Visible = on
```

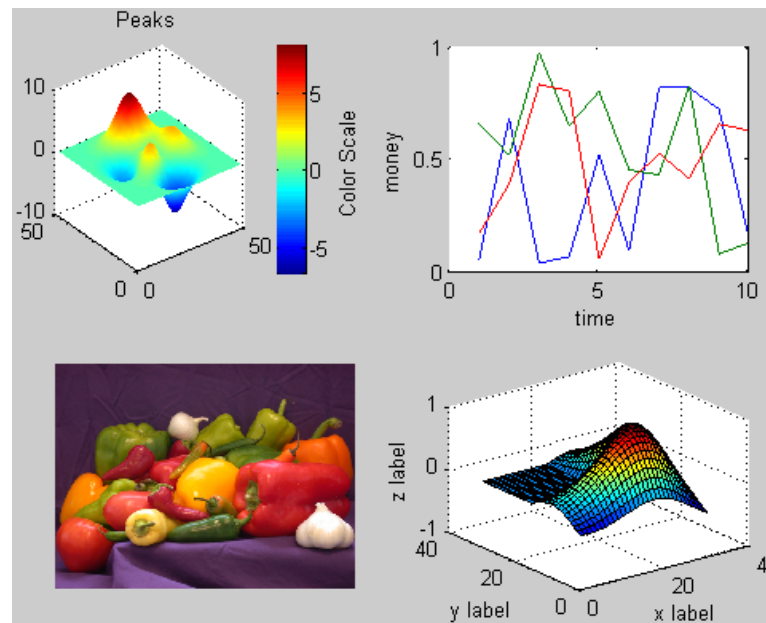
- very commonly used functions that access or modify figure properties: `figure` (creates a new figure), `clf`, `shg`, `gcf`, `tightfig`

## axes

- plot lines are children of an axes objects
- axes are children of figure objects
- figures and axes have many properties that can be set by the user
- `gca` returns a handle to the current axes
- `get(gca)` lists the attributes of an axes
- there are too many attributes to list here
- can set: fonts, font size, axis limits, grids, tick marks, tick labels, x and y labels, title, colors, viewing angle, size and position in the figure window, callback functions for user interaction, etc.
- very commonly used functions that access or modify axes properties: `axis` (on, off, square, image, equal, etc.), `xlim`, `ylim`, `grid on`, `grid off`, `cla`, `gca`, `xlabel`, `ylabel`, `title`, `hold on`, `hold off`, `tight`

# subplots

- `subplot(m,n,i)` creates an axis at position `i` in a `m` by `n` array of plots
- this is useful to place more than one plot in a figure window and on a printed page
- see also `tightfig`, `subplot_tight`, `spaceplots`, `subplotplus`, `subplot1` and `tight_subplot` (some of these available from Matlab Central)



# printing

- `print -device -options filename` prints the contents of the figure window to an image file that can be printed or included in a report
- a few of the possibilities for the `-device` flag are

```
-dps      % PostScript for black and white printers
-deps     % Encapsulated PostScript
-djpeg<nn> % JPEG image, quality level of nn (figures only)
           E.g., -djpeg90 gives a quality level of 90.
           Quality level defaults to 75 if nn is omitted.
-dtiff    % TIFF with packbits (lossless run-length encoding)
           compression (figures only)
-dtiffnocompression % TIFF without compression (figures only)
-dpng     % Portable Network Graphic 24-bit truecolor image
-dpdf     % PDF document
```

- `orient` sets the paper orientation for printing (call before `print`)

```
orient landscape;
orient rotated;
orient portrait;
orient tall;
```


## example

I used a microphone and recorded myself whistling a tune for about ten seconds and saved the data into the wave file “whistle.wav”.

# time domain plot

```
1 [x,sample_rate,nbits]=wavread('whistle.wav');
2 time = [0:length(x)-1]/sample_rate; % time [seconds]
3
4 subplot(1,2,1)
5 plot(time,x); % plot with accurately scaled time axis
6 xlabel('time [seconds]','FontSize',18);
7 ylabel('amplitude','FontSize',18);
8 set(gca,'FontSize',16)
9 grid on;
10
11 subplot(1,2,2);
12 plot(time,x); % plot with accurately scaled time axis
13 % adjust the view
14 t1 = 1.40; % seconds
15 t2 = 1.45; % seconds
16 xlim([t1 t2]); % view from 4.4 to 4.6 seconds
17 xlabel('time [seconds]','FontSize',18);
18 ylabel('amplitude','FontSize',18);
19 set(gca,'FontSize',16)
20 grid on;
21
22 orient landscape;
23 print -dpdf plotting_time.pdf
```

# time domain plot





# oscilloscope

- sometimes it is useful to visualize a signal over time

```
1 [x,fs]=audioread('.../.../dsp_labs/whistle.wav');
2 t = [0:length(x)-1]/fs; % time [seconds]
3 %plot(t,x); % this plots whole waveform, gives big picture,
4             % but cannot see detail
5
6 win_sec = 0.05; % window length [seconds]
7 win_sam = round(win_sec*fs); % window length [samples]
8 step_sec = 0.001; % step length [seconds]
9 step_sam = round(step_sec*fs); % step length [samples]
10
11 han = plot(t(1:win_sam),x(1:win_sam)); drawnow;
12 ylim(0.1*[-1 1]);
13 for i=win_sam:step_sam:length(x)
14     ind = [i-win_sam+1:i];
15     set(han,'XData',t(ind),'YData',x(ind));
16     xlim(t(ind([1,end])));
17     drawnow;
18     %pause; % if this runs too fast, include a pause
19 end
```

`image and imagesc`

## using image

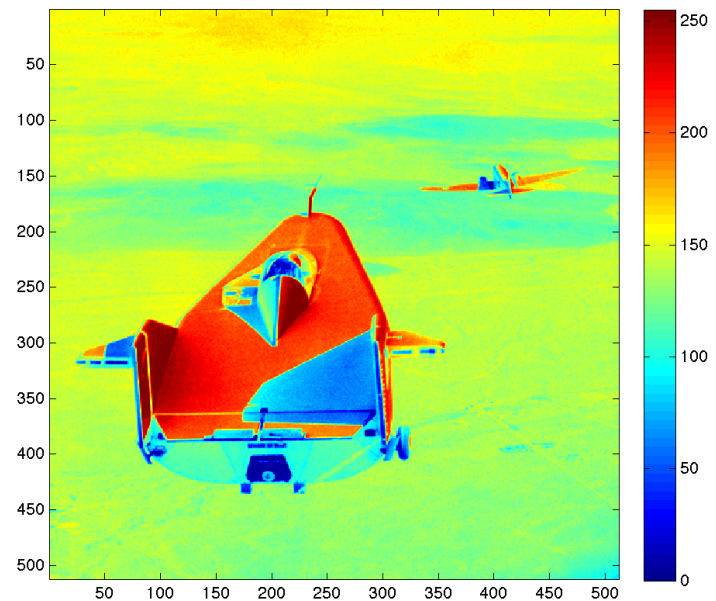
- for 2D arrays: `image(C)` displays an image using a colormap
  - `colorbar` adds a colorbar to an image
  - `colormap(gray)` changes to a grayscale colormap
  - `colormap(jet)` is the default colormap
  - built-in colormaps: `parula`, `hsv`, `hot`, `cool`, and many more
  - can also construct custom colormaps
- for 3D arrays: `image(C)` uses `C(:, :, 1)` as the red intensity, `C(:, :, 2)` as the green intensity, `C(:, :, 3)` as the blue intensity
- if `C` is double, then the elements must be in the range `[0,1]`
- if `C` is uint8, then the elements must be in the range `[0,255]`
- `image(X, Y, C)` uses the values in vectors `X` and `Y` to specify the location of the pixel centers

## example grayscale image

```
1 x = imread('liftingbody.png');  
2 h = image(x); axis image;  
3 set(h,'CDataMapping','scaled'); % I should not have to do this.  
4 colormap(jet); % set the colormap  
5 colorbar;      % show the colorbar
```



(image file)



(Matlab figure window)


## using imagesc

- `imagesc` is much easier to use – you don't have to worry about the input data range
- I almost always use `imagesc`
- `imagesc(C,[clo chi])` lets you clip the colormap

```
1 x = imread('liftingbody.png');  
2 imagesc(x,[100 200]); axis image;  
3 colormap(gray); % set the colormap  
4 colorbar;       % show the colorbar
```



(image file)



(Matlab figure window)

## imagesc with handle graphics

```
1 % load the frames of a video one at a time
2 % and display each one as an image
3
4 obj = VideoReader('xylophone.mp4','tag','myreader1');
5
6 nframes = obj.NumberOfFrames;
7 for i=1:nframes
8     y = read(obj,i); % read the next video frame
9     if(i==1)
10         han = imagesc(y); axis image; drawnow;
11     else
12         set(han,'CData',y); drawnow;
13     end
14 end
```


## FFT and spectral plots

# frequency domain plot

```
1 [x,sample_rate,nbits]=wavread('whistle.wav');
2 t1 = 1.40; % seconds
3 t2 = 1.45; % seconds
4 i1 = round(t1*sample_rate); % convert time to index
5 i2 = round(t2*sample_rate); % convert time to index
6
7 nfft = 2^12; % FFT size
8 freq = ([0:nfft-1]/nfft - 0.5)*sample_rate; % frequency [Hz]
9 X = fft(x(i1:i2),nfft); % compute the discrete-Fourier transform
10 plot(freq,20*log10(abs(fftshift(X))));
11 % plot with accurately scaled frequency axis
12
13 xlabel('frequency [Hz]','FontSize',18);
14 ylabel('magnitude [dB]','FontSize',18);
15 set(gca,'FontSize',16)
16 grid on;
17 print -dpdf plotting_freq.pdf
```



## frequency domain plot




- whistling can produce very pure tonal/sinusoidal signals
- what would this plot look like if it was a pure tone/sinusoid?

**spectrograms and time-frequency plots**

# spectrogram plot

```
1 [x,sample_rate,nbits]=wavread('whistle.wav');  
2  
3 nfft = 2^8; % FFT size  
4 overlap = round(0.8*nfft);  
5 window = hamming(nfft);  
6 spectrogram(x,window,overlap,nfft,sample_rate);  
7  
8 set(gca,'FontSize',16);  
9 grid on;  
10 print -dpdf plotting_spectrogram.pdf
```

## spectrogram (time-frequency) plot



- spectrogram shows where energy is concentrated in time and frequency

assignment

# file format

- this assignment involves reading various types of signals into Matlab and writing them out to a raw binary file
- raw binary files have a common file format regardless of the type of media
- raw binary files have two parts: a 20 byte header followed by data (in that order)
- the first 20 bytes of the file should be five integers written in the following order
  1. `ndim` - number of dimensions in the signal
  2. `nchan` - number of channels in the signal
  3. `dim0` - the size of the first dimension
  4. `dim1` - the size of the second dimension
  5. `dim2` - the size of the third dimension

(all of these integers must appear in this order regardless of the number of dimensions or channels in the signal)
- the rest of the file consists of the signal samples written as float data type
- the order in which the samples are written is important
- the order may not be the same as Matlab's natural memory layout
- some data permutations may be required

# audio file format

- audio signals are written in sample index order with samples from multiple channels interleaved
- this format enlarges the file size, but it makes processing in C easier
- for audio files, dim1 and dim2 are not used
- let dim1 be the sample rate ( $F_s$ )

ndim = 1 (int)
nchan = 2 (int)
dim0 = L (int)
dim1 = $F_s$ (int)
dim2 = 0 (int)
sample 0, chan 0 (float)
sample 0, chan 1 (float)
sample 1, chan 0 (float)
sample 1, chan 1 (float)
sample 2, chan 0 (float)
sample 2, chan 1 (float)
sample 3, chan 0 (float)
sample 3, chan 1 (float)
sample 4, chan 0 (float)
sample 4, chan 1 (float)
⋮
sample L-1, chan 0 (float)
sample L-1, chan 1 (float)

# image file format

- image signals are written in row-major order with (R,G,B) interleaved
- this format enlarges the file size, but it makes processing in C easier

ndim = 2 (int)				
nchan = 3 (int)				
dim0 = M (int)				
dim1 = N (int)				
dim2 = 0 (int)				
row 0	pixel (0,0)	R	G	B (float × 3)
	pixel (0,1)	R	G	B (float × 3)
	⋮			
	pixel (0,N-1)	R	G	B (float × 3)
row 1	pixel (1,0)	R	G	B (float × 3)
	pixel (1,1)	R	G	B (float × 3)
	⋮			
	pixel (1,N-1)	R	G	B (float × 3)
⋮	⋮			
row M-1	pixel (M-1,0)	R	G	B (float × 3)
	pixel (M-1,1)	R	G	B (float × 3)
	⋮			
	pixel (M-1,N-1)	R	G	B (float × 3)



## video file format

- the video file format follows the image format with each frame of video written out in sequence: frame 0, frame 1, ..., frame  $T-1$ , where each frame is an RGB or grayscale image
- this enlarges the file size, but it makes processing in C easier

## assignment: do the following

### 1. audio (use `flute22.wav` and `music.mp3`)

- (a) write a Matlab function named `audio2bin` that reads signal samples from an audio file (such as `.wav` or `.mp3`) and writes the signal samples to a raw binary file
- (b) write a Matlab function named `bin2audio` that reads signal samples from a raw binary file and writes the signal samples to an audio file (such as `.wav` or `.mp3`)
- (c) use the `sound` or `soundsc` Matlab commands to play an audio signal to the speaker
- (d) make a subplot plot of the audio signal using a true time-scaled x-axis (set the x-axis limits to `[1.0 1.01]` seconds)
- (e) make a subplot plot of the spectrum (FFT) of the audio signal data from 1 to 1.01 seconds and use a true frequency-scaled x-axis (what is the frequency of the highest peak?)
- (f) in a subplot show the spectrogram of the signal
- (g) write an oscilloscope script to visualize the audio signal over a 10 msec window (handle graphics)

### 2. image (use `liftingbody.png` and `coloredChips.png`)

- (a) write a Matlab function named `image2bin` that reads signal samples from an image file and writes the signal samples to a raw binary file

- (b) write a Matlab function named `bin2image` that reads signal samples from a raw binary file and writes the signal samples to an image file
- (c) make a plot of the image (pixels should be square)

### 3. video (use `xylophone.mp4`)

- (a) write a Matlab function named `video2bin` that reads signal samples from a video file and writes the signal samples to a raw binary file (only one video frame in memory at a time)
- (b) write a Matlab function named `bin2video` that reads signal samples from a raw binary file and writes the signal samples to a video file (only one video frame in memory at a time)
- (c) write a movie player script to visualize the frames of video (handle graphics)